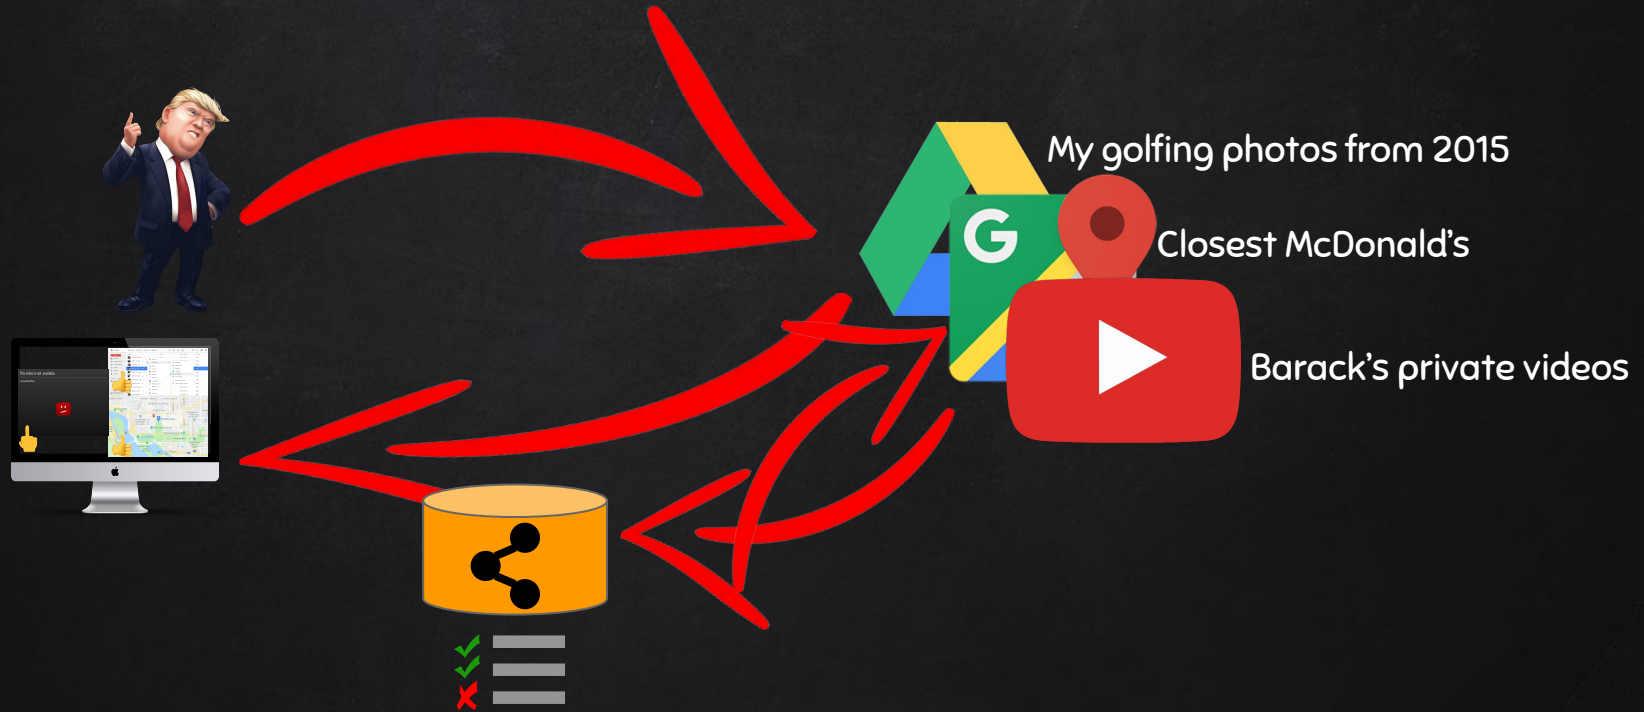




ZANZIBAR: Google's CONSISTENT, GLOBAL AUTHORIZATION SYSTEM

Presenter: Amir Ilkhechi

PRIVACY'S CORE REQUIREMENT: AUTHORIZATION CHECKS



ACLs (ACCESS CONTROL LISTS)



AUTHORIZATION ENGINE



ZANZIBAR DESIGN GOALS



Correctness

Flexibility

Scalability

Low Latency

High Availability

NAMESPACES, RELATIONS, USERSETS, AND TUPLES

Namespace: videos





Object	Relation	Userset
Video A	viewer	User A
Video B	viewer	All users
Video C	commenter	User Z

NAMESPACES, RELATIONS, USERSETS, AND TUPLES

Namespace: videos

Object	Relation	Userset
Video A	viewer	User A
Video B	viewer	All users
Video C	commenter	User Z

ACL check results:

- Video A, viewer, user A? 
- Video A, viewer, user F? 
- Video B, viewer, user B? 
- Video C, commenter, user Z? 

USERSET INDIRECTIONS...

Namespace: videos

Object	Relation	Userset
Video A	viewer	(Group G, member)
Video B	viewer	(Group K, member)
Video C	commenter	User Z

Namespace: groups

Object	Relation	Userset
Group G	member	(Group A, member)
Group K	member	User B
Group K	member	User C

deep

wide

ACL check results:

- Video B, viewer, user B? ✓
- Video B, viewer, user F? ✗

NAMESPACE CONFIG


```
name: "doc"

relation { name: "owner" }

relation {
  name: "editor"
  userset_rewrite {
    union {
      child { _this {} }
      child { computed_userset { relation: "owner" } }
    } } }

relation {
  name: "viewer"
  userset_rewrite {
    union {
      child { _this {} }
      child { computed_userset { relation: "editor" } }
      child { tuple_to_userset {
        tupleset { relation: "parent" }
        computed_userset {
          object: $TUPLE_USERSET_OBJECT # parent folder
          relation: "viewer"
        }
      } } }
    } } }
} } }
```

"NEW ENEMY" PROTECTION 1



“NEW ENEMY” PROTECTION 2




CONSISTENCY PROTOCOL

Updates by Alice:



updateACL(doc x, viewer, remove Bob)

Timestamp T_0




Content update by
Charlie:



CheckContentUpdate(doc x, writer, Charlie)

Yes, Timestamp T_1 [$T_1 > T_0$]





ACL check for Bob:

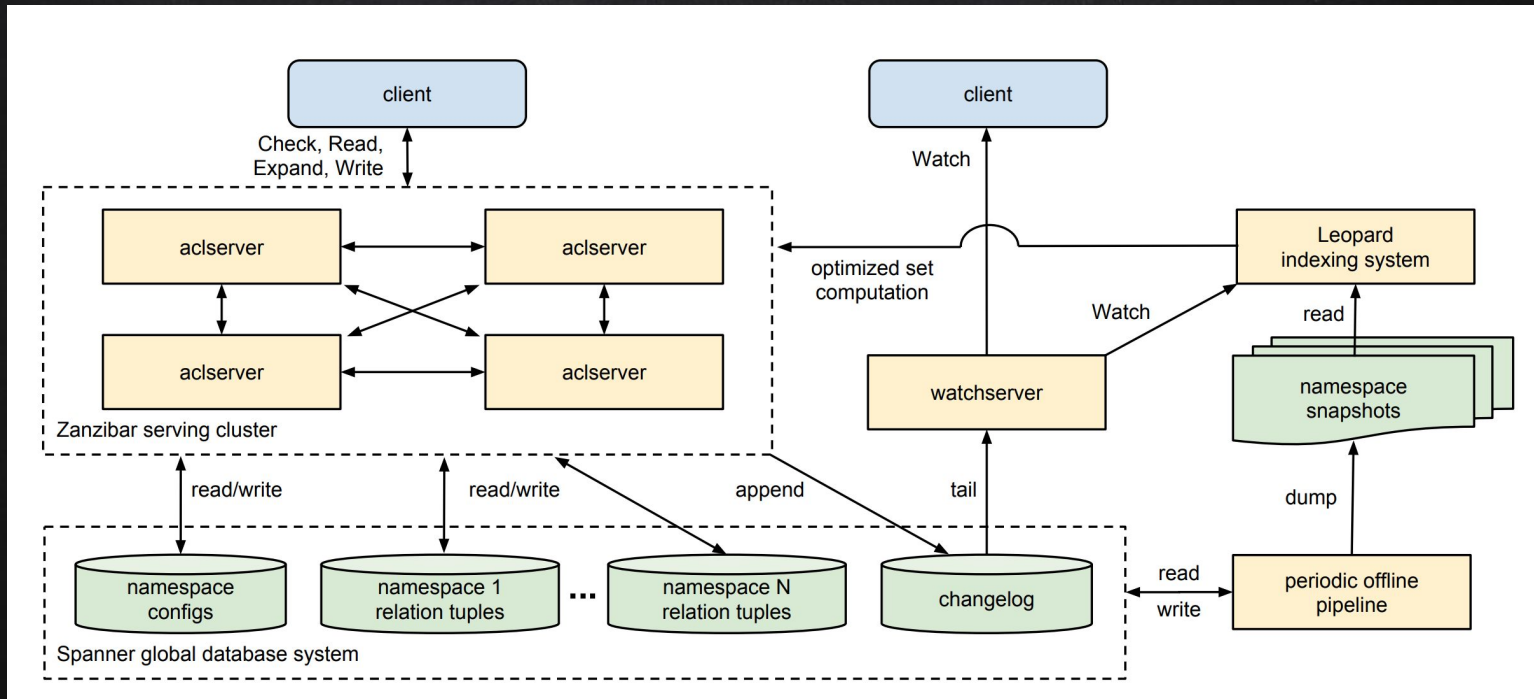


CheckACL(doc x, viewer, Bob, T_1)

No [at $T_2 > T_1 > T_0$]



ARCHITECTURE




OPTIMIZATIONS

- Timestamps are chosen to reduce latency
- Hot-spot mitigation
- Request hedging
- Isolation
- Optimized processing of nested sets


ZANZIBAR DEPLOYMENT

- In production use for 5+ years
- Has 1500+ namespaces
- 2+ trillion relation tuples
- 10+ million queries per second
- 10+ thousand servers


RATE OF CHECK QUERIES



CHECK SAFE LATENCY



AVAILABILITY



SUMMARY

- Robust authorization checks → essential to preserving privacy
- Zanzibar offers a unified authorization system:
 - ◆ Causal ordering of actions
 - ◆ Rich spectrum of access control policies
 - ◆ Low latency and high availability
 - ◆ Scalable
 - ◆ Flexible