

CryptDB

- Pessimistic - Data gets stolen all the time
 - 1: DBAs can look but should not ("curious DBA")
 - "Malicious DBAs are more like to read data... than to change data or query results"
 - 2: Adversary gains control (General theft)
 - Not trusted: DBMS machines, administrators
 - Trusted: application and proxy
- Application <-> Proxy <-> Server
 - Good at doing predefined queries. Good? Bad
 - Proxy intercepts query from Application, encrypts, and sends to server
 - Server does its thing, sends to proxy, proxy decrypts, sends to app
 - Master key in the proxy
 - All on the same machine according to MR
 - DBMS sees anonymized schema (puzzle pieces with no picture)
- **Idea: execute SQL queries over encrypted data**
- We could use randomized encryption
 - Q: What happens if we ask for people with Employee ID of 1234?
 - A: Could map to different values
 - S: Use non-deterministic encryption for equality checks
 - Server can perform equality checks by checking the encrypted data
- Good news: most of SQL is just a handful of operations
- Equality
- Note that the ORDER IS NOT PRESERVED
 - Use an order-preserving encryption (OPE) scheme
- Encryption schemes (decreasing security, increasing functionality):
 - Randomized (RAND)
 - strongest, leaks no info about the data
 - No SQL use
 - Homomorphic (HOM)
 - SUM operations
 - $HOM(x) \cdot HOM(y) = HOM(x+y)$
 - pretty secure (almost RAND)
 - Stan Zdonik
 - Search (SEARCH)
 - LIKE operations
 - Note: only full words supported
 - Deterministic (DET)
 - Equality operations (=, in, groups)
 - Join (JOIN)
 - Finding equality matches between two columns
 - Q: Each column has a different key!
 - Order Preserving (OPE)
 - $x < y \Rightarrow OPE(x) < OPE(y)$
- **Idea: use "onions" to encrypt in layers ("adjustable query-based encryption")**
 - We NEVER fully decrypt an onion; there is ALWAYS at least one layer
 - Eq
 - Ord
 - Search
 - Add
 - Note that the RAND is on the outside!
 - We encrypt from weakest on inside to strongest on outside
 - Q: Why multiple onions?
 - A: Computations supported by different enc schemes not necc. ord'd
 - A: Performance considerations
 - Problems? Space concerns? How much concern?
 - Where does a specific onion make sense?
 - Search on integers? NO
 - Search on username? YES
 - Add on strings? NO
- ALL keys are derived from the MK
 - Each key is specific to a single COL

- User keys are generated at each login, deleted at log off
- Onion decryption is performed entirely by the DBMS server. Why?
 - Only done when needed / DBMS is secure
- When you are offline, your key is not available
 - What if I send a message to Alice while she's offline?
 - Encrypts message with Alice's public key, Alice logs on and automatically encrypts with her private key
- NO change to the DBMS, all through UDFs. Good? Bad?
- Not all operators supported
 - Trigonometry
 - Combining encryption schemes ($T1.a + T1.b > T2.c$)
 - HOM does not preserve order
 - Is this a problem?
- How do we set this up:
 - Step 1: Define principle types (external: authenticated users, internal: delegated users)
 - Step 2: Specify sensitive data columns ENC_FOR
- Problems:
 - does NOT ensure integrity, freshness, or completeness of results
 - "Equality checks for a few rows... require revealing that class of computation for an entire column"
- How do we attack this?
 - Sorting attack: Attack OPE columns, every value sorted in the column, map to plain text
 - DBMS uses sorted intermediate results for answering queries with range joins
- What if we attack users?
 - Only logged in users are vulnerable
 - How many people stay logged into an application?
 - Case of Marriott -> most people would not be affected
 - Task: get somebody online
 - Who do we target?
- Attack: this is easily crackable because OPE is bad!
 - Response: you're not using it properly!

EX:

| GuestId | Name | Username | Email | Password | CreditCardNo

Connor stayed at the Westin and wants to pay:

SELECT CreditCardNo FROM Guests WHERE Name='Connor Luckett';

This is targeting an Eq onion. We use:

K_{Guests, CreditCardNo, Eq, RND}

K_{Guests, CreditCardNo, Eq, DET}

K_{Guests, CreditCardNo, Eq, JOIN}

Connor moved and needs to check his bill:

CREATE TABLE bills (...)

payerId int, *total* int

(*payerId* payer) int SPEAKS_FOR (*total* bill)

)